

Programming ex1

1. A simple MATLAB function


```
1 A = eye(5);
```

```
ans = 5x5
     1     0     0     0     0
     0     1     0     0     0
     0     0     1     0     0
     0     0     0     1     0
     0     0     0     0     1
```

2. Linear regression with one variable

2.1 Plotting the data

```
1 data = load('ex1data1.txt'); % read comma separated data
2 X = data(:, 1); y = data(:, 2);
3 plot(x, y, 'rx', 'MarkerSize', 10); % Plot the data
4 ylabel('Profit in $10,000s'); % Set the y-axis label
5 xlabel('Population of City in 10,000s'); % Set the x-axis label
```



2.2 Gradient Descent

2.2.1 Update Equations

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$h_{\theta}(x) = \theta^T x = \theta_0 + \theta_1 x_1$$

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \quad (\text{simultaneously update } \theta_j \text{ for all } j)$$

2.2.2 Implementation

```
1 m = length(X); % number of training examples
2 X = [ones(m,1),data(:,1)]; % Add a column of ones to x
```

```

3 theta = zeros(2, 1); % initialize fitting parameters
4 iterations = 1500;
5 alpha = 0.01;

```

2.2.3 Computing the cost $J(\theta)$

```

1 function J = computeCost(X, y, theta)
2 %COMPUTECOST Compute cost for linear regression
3 % J = COMPUTECOST(X, y, theta) computes the cost of using theta as the
4 % parameter for linear regression to fit the data points in X and y
5 % Initialize some useful values
6 m = length(y); % number of training examples
7 % You need to return the following variables correctly
8 J = 0;
9 % ===== YOUR CODE HERE =====
10 % Instructions: Compute the cost of a particular choice of theta
11 % You should set J to the cost.
12 J = 1/(2*m) * sum((X * theta - y).^2);
13 % =====
14 end
15
16 % Compute and display initial cost with theta all zeros
17 computeCost(X, y, theta)

```

ans = 32.0727

```

1 % Compute and display initial cost with non-zero theta
2 computeCost(X, y, [-1; 2])

```

ans = 54.2425

2.2.4 Gradient descent

```

1 function [theta, J_history] = gradientDescent(X, y, theta, alpha, num_iters)
2 %GRADIENDESCENT Performs gradient descent to learn theta
3 % theta = GRADIENDESCENT(X, y, theta, alpha, num_iters) updates theta by
4 % taking num_iters gradient steps with learning rate alpha
5 % Initialize some useful values
6 m = length(y); % number of training examples
7 J_history = zeros(num_iters, 1);
8 for iter = 1:num_iters
9     % ===== YOUR CODE HERE =====
10    % Instructions: Perform a single gradient step on the parameter vector
11    % theta.
12    %
13    % Hint: While debugging, it can be useful to print out the values
14    % of the cost function (computeCost) and gradient here.

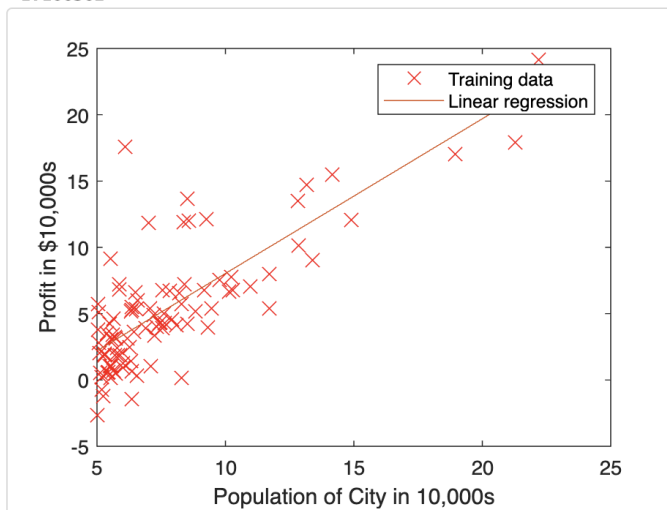
```

```

15 %
16 theta = theta - (alpha/m) * (X' * (X * theta - y));
17 % =====
18 % Save the cost J in every iteration
19 J_history(iter) = computeCost(X, y, theta);
20 end
21 end
22
23 % Run gradient descent:
24 % Compute theta
25 theta = gradientDescent(X, y, theta, alpha, iterations);
26
27 % Print theta to screen
28 % Display gradient descent's result
29 fprintf('Theta computed from gradient descent:\n%f,\n%f',theta(1),theta(2))
30
31 % Plot the linear fit
32 hold on; % keep previous plot visible
33 plot(X(:,2), X*theta, '-')
34 legend('Training data', 'Linear regression')
35 hold off % don't overlay any more plots on this figure
36
37 % Predict values for population sizes of 35,000 and 70,000
38 predict1 = [1, 3.5] * theta;
39 fprintf('For population = 35,000, we predict a profit of %f\n', predict1*1000)
40 predict2 = [1, 7] * theta;
41 fprintf('For population = 70,000, we predict a profit of %f\n', predict2*1000)

```

Theta computed from gradient descent:
-3.630291,
1.166362



For population = 35,000, we predict a profit of 4519.767868
For population = 70,000, we predict a profit of 45342.450129

2.4 Visualizing $J(\theta)$

```


1 % Visualizing J(theta_0, theta_1):
2 % Grid over which we will calculate J
3 theta0_vals = linspace(-10, 10, 100);
4 theta1_vals = linspace(-1, 4, 100);

```

```

5
6 % initialize J_vals to a matrix of 0's
7 J_vals = zeros(length(theta0_vals), length(theta1_vals));
8
9 % Fill out J_vals
10 for i = 1:length(theta0_vals)
11     for j = 1:length(theta1_vals)
12         t = [theta0_vals(i); theta1_vals(j)];
13         J_vals(i,j) = computeCost(X, y, t);
14     end
15 end
16
17 % Because of the way meshgrids work in the surf command, we need to
18 % transpose J_vals before calling surf, or else the axes will be flipped
19 J_vals = J_vals';
20
21 % Surface plot
22 figure;
23 surf(theta0_vals, theta1_vals, J_vals)
24 xlabel('\theta_0'); ylabel('\theta_1');
25
26 % Contour plot
27 figure;
28 % Plot J_vals as 15 contours spaced logarithmically between 0.01 and 100
29 contour(theta0_vals, theta1_vals, J_vals, logspace(-2, 3, 20))
30 xlabel('\theta_0'); ylabel('\theta_1');
31 hold on;
32 plot(theta(1), theta(2), 'rx', 'MarkerSize', 10, 'LineWidth', 2);
33 hold off;
34
35

```



3. Linear regression with multiple variables

```

1 % Load Data
2 data = load('ex1data2.txt');
3 X = data(:, 1:2);
4 y = data(:, 3);
5 m = length(y);
6
7 % Print out some data points
8 % First 10 examples from the dataset

```

```
9 fprintf(' x = [%0f %0f], y = %0f \n', [X(1:10,:) y(1:10,:)]);
```

```
x = [2104 3], y = 399900
x = [1600 3], y = 329900
x = [2400 3], y = 369000
x = [1416 2], y = 232000
x = [3000 4], y = 539900
x = [1985 4], y = 299900
x = [1534 3], y = 314900
x = [1427 3], y = 198999
x = [1380 3], y = 212000
x = [1494 3], y = 242500
```

3.1 Feature Normalization

```
1 function [X_norm, mu, sigma] = featureNormalize(X)
2 %FEATURENORMALIZE Normalizes the features in X
3 %   FEATURENORMALIZE(X) returns a normalized version of X where
4 %   the mean value of each feature is 0 and the standard deviation
5 %   is 1. This is often a good preprocessing step to do when
6 %   working with learning algorithms.
7 % You need to set these values correctly
8 X_norm = X;
9 mu = zeros(1, size(X, 2));
10 sigma = zeros(1, size(X, 2));
11 % ===== YOUR CODE HERE =====
12 % Instructions: First, for each feature dimension, compute the mean
13 %               of the feature and subtract it from the dataset,
14 %               storing the mean value in mu. Next, compute the
15 %               standard deviation of each feature and divide
16 %               each feature by it's standard deviation, storing
17 %               the standard deviation in sigma.
18 %
19 %               Note that X is a matrix where each column is a
20 %               feature and each row is an example. You need
21 %               to perform the normalization separately for
22 %               each feature.
23 %
24 % Hint: You might find the 'mean' and 'std' functions useful.
25 %
26 mu = mean(X_norm);
27 sigma = std(X_norm);
28 for i = 1 : size(X, 2)
29     X_norm(:, i) = X_norm(:, i) - mu(1, i);
30     X_norm(:, i) = X_norm(:, i) ./ sigma(1, i);
31 end
32 % =====
33 end
```

```
1 % Scale features and set them to zero mean
2 [X, mu, sigma] = featureNormalize(X);
```

```

3 % Add intercept term to X
4 X = [ones(m, 1) X];

```

3.2 Gradient Descent

$$J(\theta) = \frac{1}{2m} (X\theta - \vec{y})^T (X\theta - \vec{y})$$

$$X = \begin{bmatrix} - (x^{(1)})^T & - \\ - (x^{(2)})^T & - \\ \vdots & \\ - (x^{(m)})^T & - \end{bmatrix} \quad \vec{y} = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{bmatrix}$$

```

1 function [theta, J_history] = gradientDescentMulti(X, y, theta, alpha, num_iters)
2 %GRADIENDESCENTMULTI Performs gradient descent to learn theta
3 %   theta = GRADIENDESCENTMULTI(x, y, theta, alpha, num_iters) updates theta by
4 %   taking num_iters gradient steps with learning rate alpha
5 % Initialize some useful values
6 m = length(y); % number of training examples
7 J_history = zeros(num_iters, 1);
8 for iter = 1:num_iters
9     % ===== YOUR CODE HERE =====
10    % Instructions: Perform a single gradient step on the parameter vector
11    %                  theta.
12    %
13    % Hint: While debugging, it can be useful to print out the values
14    %        of the cost function (computeCostMulti) and gradient here.
15    %
16    theta = theta - (alpha/m) * (X' * (X * theta - y));
17    % =====
18    % Save the cost J in every iteration
19    J_history(iter) = computeCostMulti(X, y, theta);
20 end
21 end

```

```

1 function J = computeCostMulti(X, y, theta)
2 %COMPUTECOSTMULTI Compute cost for linear regression with multiple variables
3 %   J = COMPUTECOSTMULTI(X, y, theta) computes the cost of using theta as the
4 %   parameter for linear regression to fit the data points in X and y
5 % Initialize some useful values
6 m = length(y); % number of training examples
7 % You need to return the following variables correctly
8 J = 0;
9 % ===== YOUR CODE HERE =====
10 % Instructions: Compute the cost of a particular choice of theta
11 %               You should set J to the cost.
12 J = 1/(2*m) * sum((X * theta - y) .^2);
13 % =====

```

14 end

```
1 % Run gradient descent
2 % Choose some alpha value
3 alpha = 0.1;
4 num_iters = 400;
5
6 % Init Theta and Run Gradient Descent
7 theta = zeros(3, 1);
8 [theta, ~] = gradientDescentMulti(X, y, theta, alpha, num_iters);
9
10 % Display gradient descent's result
11 fprintf('Theta computed from gradient descent:\n%f\n%f\n%f', theta(1), theta(2), theta(3))
```

Theta computed from gradient descent:
340412.659574
110631.048958
-6649.472950


```
1 % Estimate the price of a 1650 sq-ft, 3 br house
2 % ===== YOUR CODE HERE =====
3
4 param = [1650, 3];
5 norm_param = (param - mu) ./ sigma;
6 real_param = [1, norm_param];
7
8
9 price = real_param * theta; % Enter your price formula here
10
11 % =====
12
13 fprintf('Predicted price of a 1650 sq-ft, 3 br house (using gradient descent) = $%.2f\n', price)
```

Predicted price of a 1650 sq-ft, 3 br house (using gradient descent):
\$293081.464622

3.2.1 Optional (ungraded) exercise: Selecting learning rates

```
1 % Run gradient descent:
2 % Choose some alpha value
3 alpha = 0.03;
4 num_iters = 200;
5
6 % Init Theta and Run Gradient Descent
7 theta = zeros(3, 1);
8 [~, J_history] = gradientDescentMulti(X, y, theta, alpha, num_iters);
9
10 % Plot the convergence graph
11 plot(1:num_iters, J_history, '-b', 'LineWidth', 2);
12 xlabel('Number of iterations');
```

```
13 ylabel('Cost J');
```



```
1 % Run gradient descent
2 % Replace the value of alpha below best alpha value you found above
3 alpha = 0.03;
4 num_iters = 200;
5
6 % Init Theta and Run Gradient Descent
7 theta = zeros(3, 1);
8 [theta, ~] = gradientDescentMulti(X, y, theta, alpha, num_iters);
9
10 % Display gradient descent's result
11 fprintf('Theta computed from gradient descent:\n%f\n%f\n%f', theta(1), theta(2), theta(3));
12
13 % Estimate the price of a 1650 sq-ft, 3 br house. You can use the same
14 % code you entered earlier to predict the price
15 % ===== YOUR CODE HERE =====
16 param = [1650, 3];
17 norm_param = (param - mu) ./ sigma;
18 real_param = [1, norm_param];
19
20
21 price = real_param * theta;
22 % =====
23
24 fprintf('Predicted price of a 1650 sq-ft, 3 br house (using gradient descent)');
```

```
Theta computed from gradient descent:
339642.904508
106274.973805
-2302.198941
```

```
Predicted price of a 1650 sq-ft, 3 br house (using gradient descent):
$293261.551453
```

3.3 Normal Equations

$$\theta = (X^T X)^{-1} X^T \vec{y}$$

```
1 function [theta] = normalEqn(X, y)
2 %NORMALEQN Computes the closed-form solution to linear regression
3 % NORMALEQN(X,y) computes the closed-form solution to linear
4 % regression using the normal equations.
5 theta = zeros(size(X, 2), 1);
```



```

6 % ===== YOUR CODE HERE =====
7 % Instructions: Complete the code to compute the closed form solution
8 %           to linear regression and put the result in theta.
9 %
10 % ----- Sample Solution -----
11 theta = pinv(X' * X) * X' * y;
12 % -----
13 % =====
14 end

```

```

1 % Solve with normal equations:
2 % Load Data
3 data = csvread('ex1data2.txt');
4 X = data(:, 1:2);
5 y = data(:, 3);
6 m = length(y);
7
8 % Add intercept term to X
9 X = [ones(m, 1) X];
10
11 % Calculate the parameters from the normal equation
12 theta = normalEqn(X, y);
13
14 % Display normal equation's result
15 fprintf('Theta computed from the normal equations:\n%f\n%f\n%f', theta(1), theta(2), theta(3));

```

```

Theta computed from the normal equations:
89597.909544
139.210674
-8738.019113

```

```

1 % Estimate the price of a 1650 sq-ft, 3 br house.
2 % ===== YOUR CODE HERE =====
3 param = [1650, 3];
4 real_param = [1, param];
5 price = real_param * theta; % Enter your price formula here
6
7 % =====
8
9 fprintf('Predicted price of a 1650 sq-ft, 3 br house (using normal equations): $%f', price);

```

```

Predicted price of a 1650 sq-ft, 3 br house (using normal equations):
$293081.464335

```

